

Programming for BIG Data

11. Scientific Python

11. Scientific Python



lpython

NumPy

SciPy

ScikitLearn

Matplotlib

Pandas

The ideal programming environment for computational mathematics enjoys the following characteristics:

- It must be based on a computer language that allows the user to work quickly and integrate systems effectively. Ideally, the computer language should be portable to all platforms: Windows, Mac OS X, Linux, Unix, Android, and so on.
- Besides running the compiled code, the programming environment should allow the possibility of interactive sessions as well as scripting capabilities for quick experimentation.
- Different coding paradigms should be supported—imperative, object-oriented, and/or functional coding styles.
- It should be an open source software, that allows user access to the raw data code, and allows the user to modify basic algorithms if so desired.

What's out there ...

Among the best-known environments for numerical computations used by the scientific community is **MATLAB**, which is commercial, expensive, and which does not allow any tampering with the code.

Maple and **Mathematica** are more geared towards symbolic computation, although they can match many of the numerical computations from MATLAB. These are, however, also commercial, expensive, and closed to modifications.

One environment that combines the best of all worlds is Python with the open source libraries **NumPy** and **SciPy** for numerical operations.

The first property that attracts users to Python is, without a doubt, its code readability. The syntax is extremely clear and expressive.

It has the advantage of supporting code written in different paradigms: object oriented, functional, or old school imperative.

Python comes equipped with a large library of packages for machine learning tasks:

Python comes equipped with a large library of packages for machine learning tasks:

The **IPython** console.



NumPy, which is an extension that adds support for multi-dimensional arrays, matrices, and high-level mathematical functions.



SciPy, which is a library of scientific formulae, constants, and mathematical functions.



Scikit-learn, which is a library for machine learning tasks
Such as classification, regression, and clustering



Python comes equipped with a large library of packages for machine learning tasks:

Matplotlib, which is for creating plots



Pandas provides fast, flexible
and expressive data structures designed
to make working with structured



(tabular, multidimensional, potentially heterogeneous) and time series data both
easy and intuitive. It aims to be the fundamental high-level building block for
doing practical, real world data analysis in Python.

The IPython console.

At the core of IPython is the IPython console: a powerful interactive interpreter that allows you to test your ideas in a very fast and intuitive way. Instead of having to create, save, and run a file every time you want to test a code snippet, you can simply type it into a console.

A powerful feature of IPython is that it decouples the traditional read-evaluate-print loop that most computing platforms are based on.

IPython puts the evaluate phase into its own process: a kernel (not to be confused with the kernel function used in machine learning algorithms).

More than one client can access the kernel. This means you can run code in a number of files and access them, for example, running a method from the console. Also, the kernel and the client do not need to be on the same machine. This has powerful implications for distributed and networked computing.



```
IPython: C:\Users\Jake
Microsoft Windows [Version 10.0.14393]
(c) 2016 Microsoft Corporation. All rights reserved.

C:\Users\Jake>ipython
Python 3.5.2 [Anaconda custom (64-bit)] (default, Jul 5 2016, 11:41:13)
Type "copyright", "credits" or "license" for more information.

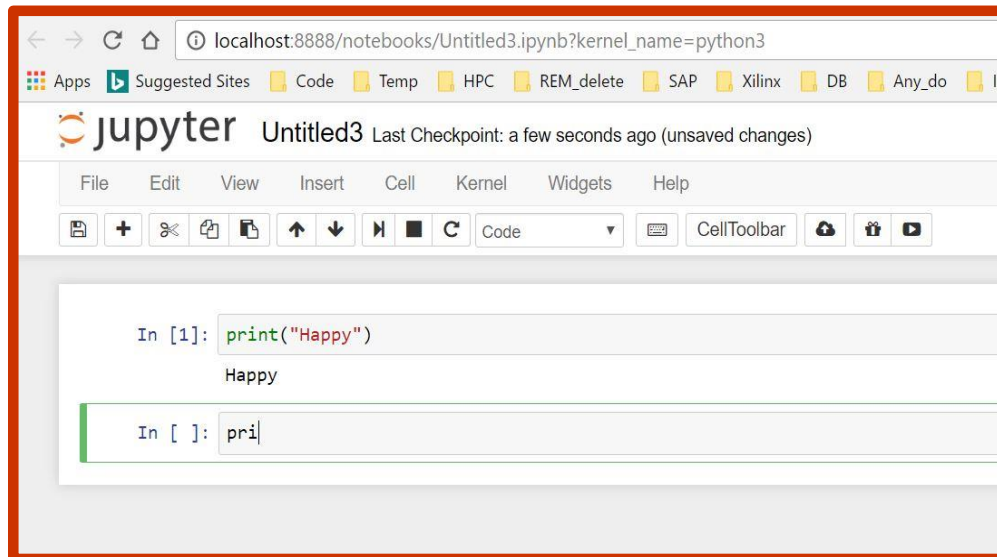
IPython 5.1.0 -- An enhanced Interactive Python.
?      -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help    -> Python's own help system.
object? -> Details about 'object', use 'object??' for extra details.

In [1]: print("Happy")
Happy

In [2]: pri
```



The IPython notebook has merged into another project known as Jupyter (jupyter.org).



This web application is a powerful platform for numerical computing in over 40 languages.

The notebook allows you to share and collaborate on live code and publish rich graphics and text.

There is a hierarchy of types for representing data in Python.

At the root are **immutable** objects such as integers, floats, and Boolean.

Built on this, we have sequence types. These are ordered sets of objects indexed by non-negative integers. They are iterative objects that include strings, lists, and tuples. Sequence types have a common set of operations such as returning an element (`s[i]`) or a slice (`s[i : j]`), and finding the length (`len(s)`) or the sum (`sum(s)`).

Finally, we have mapping types. These are collections of objects indexed by another collection of key objects. Mapping objects are unordered and are indexed by numbers, strings, or other objects. The built-in Python mapping type is the dictionary.

NumPy builds on these data objects by providing two further objects: an N-dimensional array object (**ndarray**) and a universal function object (**ufunc**).

The ufunc object provides element-by-element operations on ndarray objects, allowing typecasting and array broadcasting. Typecasting is the process of changing one data type into another, and broadcasting describes how arrays of different sizes are treated during arithmetic operations.

There are sub-packages for:

- Linear algebra (linalg)
- Random number generation (random)
- Discrete Fourier transforms (fft)
- Unit testing (testing)

Indexing and slicing in NumPy builds on the slicing and indexing techniques used in sequences.

We are already familiar with slicing sequences, such as lists and tuples, in Python using the `[i : j : k]` syntax, where:

`i` is the start index,

`j` is the end,

`k` is the step.

NumPy extends this concept of the selection tuple to N-dimensions.

Lets try some code



```
numpy.arange([start, ]stop, [step, ]dtype=None)
```

Return evenly spaced values within a given interval.

Values are generated within the half-open interval `[start, stop)` (in other words, the interval including start but excluding stop). For integer arguments the function is equivalent to the Python built-in `range` function, but returns an `ndarray` rather than a list.

```
>>> np.arange(3)
array([0, 1, 2])
>>> np.arange(3.0)
array([ 0.,  1.,  2.])
>>> np.arange(3,7)
array([3, 4, 5, 6])
>>> np.arange(3,7,2)
array([3, 5])
```

Lets try some code



Lets create an array of 60 integers

```
In [1]: import numpy as np
```

```
In [2]: a = np.arange(60)
```

```
In [3]: print(a)
```

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24
 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49
 50 51 52 53 54 55 56 57 58 59]
```

Lets slice it... a[???



```
In [23]: a
```

```
Out[23]:
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
        34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
        51, 52, 53, 54, 55, 56, 57, 58, 59])
```

```
In [24]: a[6]
```

```
Out[24]: 6
```

```
In [25]: a[2:5]
```

```
Out[25]: array([2, 3, 4])
```

```
In [26]: a[2:12:2]
```

```
Out[26]: array([ 2,  4,  6,  8, 10])
```

```
In [27]: a[-1]
```

```
Out[27]: 59
```

```
In [28]: a[-10:-1]
```

```
Out[28]: array([50, 51, 52, 53, 54, 55, 56, 57, 58])
```

Lets try some code



`numpy.ndarray.shape`

May be used to “reshape” the array, as long as this would not require a change in the total number of elements

```
In [4]: a.reshape(5,12)
```

```
Out[4]:
```

```
array([[ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11],
       [12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23],
       [24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35],
       [36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47],
       [48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59]])
```


Lets try some code



`numpy.ndarray.shape`

May be used to “reshape” the array, as long as this would not require a change in the total number of elements

```
In [4]: a.reshape(5,12)
```

```
Out[4]:
```

```
array([[ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11],
       [12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23],
       [24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35],
       [36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47],
       [48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59]])
```

Lets try some code



`numpy.ndarray.shape`

May be used to “reshape” the array, as long as this would not require a change in the total number of elements.

Lets re shape out array into a 5 by 12 structure → `a.reshape(?,?)`

```
In [4]: a.reshape(5,12)
```

```
Out[4]:
```

```
array([[ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11],
       [12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23],
       [24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35],
       [36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47],
       [48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59]])
```

```
In [11]: a = a.reshape(3,4,5)
```

```
In [12]: a
```

```
Out[12]:
```

```
array([[[ 0,  1,  2,  3,  4],
        [ 5,  6,  7,  8,  9],
        [10, 11, 12, 13, 14],
        [15, 16, 17, 18, 19]],

       [[20, 21, 22, 23, 24],
        [25, 26, 27, 28, 29],
        [30, 31, 32, 33, 34],
        [35, 36, 37, 38, 39]],

       [[40, 41, 42, 43, 44],
        [45, 46, 47, 48, 49],
        [50, 51, 52, 53, 54],
        [55, 56, 57, 58, 59]]])
```

```
In [14]: a[2,1,3]
Out[14]: 48
```

Lets slice it... `a[???`



Using the ellipse (...), we can select any remaining unspecified dimensions. For example, `a[..., I]` is equivalent to `a[:, :, I]`

What will `a[:, :, I]` give us ?

Lets slice it... `a[??]`



With slicing, we are creating views; the original array remains untouched, and the view retains a reference to the original array. This means that when we create a slice, even though we assign it to a new variable, if we change the original array, these changes are also reflected in the new array.

```
In [40]: b = a[2,2,0:2]
```

```
In [41]: b  
Out[41]: array([50, 51])
```

```
In [42]: print(b)  
[50 51]
```

```
In [43]: a[2] = 0
```

```
In [44]: a  
Out[44]:  
array([[ 0,  1,  2,  3,  4],  
       [ 5,  6,  7,  8,  9],  
       [10, 11, 12, 13, 14],  
       [15, 16, 17, 18, 19]],  
        
      [[20, 21, 22, 23, 24],  
       [25, 26, 27, 28, 29],  
       [30, 31, 32, 33, 34],  
       [35, 36, 37, 38, 39]],  
        
      [[ 0,  0,  0,  0,  0],  
       [ 0,  0,  0,  0,  0],  
       [ 0,  0,  0,  0,  0],  
       [ 0,  0,  0,  0,  0]])
```

Copies deep / shallow



Here, *a* and *b* are referring to the same array. When we assign values in *a*, this is also reflected in *b*. To copy an array rather than simply make a reference to it, we use the `deep copy()` function from the `copy` package in the standard library:

```
import copy  
c = copy.deepcopy(a)
```

Here, we have created a new independent array, *c*. Any changes made in array *a* will **not** be reflected in array *c*.

Lets slice it... a[???



```
In [45]: b
Out[45]: array([0, 0])
```

```
In [46]:
```

```
In [46]: import copy
```

```
In [47]: c = copy.deepcopy(a)
```

```
In [48]: c
Out[48]:
array([[ 0,  1,  2,  3,  4],
       [ 5,  6,  7,  8,  9],
       [10, 11, 12, 13, 14],
       [15, 16, 17, 18, 19]],

      [[20, 21, 22, 23, 24],
       [25, 26, 27, 28, 29],
       [30, 31, 32, 33, 34],
       [35, 36, 37, 38, 39]],

      [[ 0,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0]])
```

```
In [49]: c[1]
Out[49]:
array([[20, 21, 22, 23, 24],
       [25, 26, 27, 28, 29],
       [30, 31, 32, 33, 34],
       [35, 36, 37, 38, 39]])
```

```
In [50]: a[1] = 0
```

```
In [51]: a[1]
Out[51]:
array([[0, 0, 0, 0, 0],
       [0, 0, 0, 0, 0],
       [0, 0, 0, 0, 0],
       [0, 0, 0, 0, 0]])
```

```
In [52]: c[1]
Out[52]:
array([[20, 21, 22, 23, 24],
       [25, 26, 27, 28, 29],
       [30, 31, 32, 33, 34],
       [35, 36, 37, 38, 39]])
```

Numpy support vectorized operations



```
In [54]: a = np.arange(5)
```

```
In [55]: a
```

```
Out[55]: array([0, 1, 2, 3, 4])
```

```
In [56]: a * 2
```

```
Out[56]: array([0, 2, 4, 6, 8])
```

```
In [57]:
```


As you would expect, you can perform mathematical operations such as addition, subtraction, multiplication, as well as the trigonometric functions on NumPy arrays.

Arithmetic operations on different shaped arrays can be carried out by a process known as **broadcasting**. When operating on two arrays, NumPy compares their shapes element-wise from the trailing dimension. Two dimensions are compatible if they are the same size, or if one of them is 1. If these conditions are not met, then a `ValueError` exception is thrown.

This is all done in the background using the ufunc object. This object operates on ndarrays on a element-by-element basis. They are essentially wrappers that provide a consistent interface to scalar functions to allow them to work with NumPy arrays. There are over 60 ufunc objects covering a wide variety of operations and types. The ufunc objects are called automatically when you perform operations such as adding two arrays using the `+` operator


Matplotlib is an essential tool for visualizing data in Python.

It has a sub package PyPlot with a range of graphical functions that can be invoked on a PyPlot instance. At the core of PyPlot is the plot method. The simplest implementation is to pass plot a list or a 1D array.

If only one argument is passed to plot, it assumes it is a sequence of y values, and it will automatically generate the x values.

More commonly, we pass plot two 1D arrays or lists for the co-ordinates x and y. The plot method can also accept an argument to indicate line properties such as line width, color, and style.

```
5 @author: Jake
6 """
7
8 import numpy as np
9 import matplotlib.pyplot as plt
10
11 x = np.arange(0.,5., 0.2)
12 |
13 plt.plot(x, x**4, 'red', x, x*90, \
14          'blue', x, x**3, 'green')
15
16 plt.show()
```



This code prints three lines in different styles: a red line, blue squares, and green triangles. Notice that we can pass more than one pair of coordinate arrays to plot multiple lines. For a full list of line styles, type the `help(plt.plot)` function.


Pyplot applies plotting commands to the current axes. Multiple axes can be created using the subplot command.

The subplot() command specifies numrows, numcols, fignum where fignum ranges from 1 to numrows*numcols.

```

0
7
8 import numpy as np
9 import matplotlib.pyplot as plt
10
11 x1 = np.arange(0., 5., 0.2)
12 x2 = np.arange(0., 5., 0.1)
13
14 plt.figure(1)
15 plt.subplot(211)
16
17 plt.plot(x1, x1**4, 'r', x1, x1*90, \
18         'bs', x1, x1**3, 'g^', \
19         linewidth=2.0)
20 |
21 plt.subplot(212)
22
23 plt.plot(x2, np.cos(2*np.pi*x2), 'k')
24 plt.show()

```




Another useful plot is the histogram. The `hist()` object takes an array, or a sequence of arrays, of input values. The second parameter is the number of bins.

Lets create a histogram with labelled axes, and use 10 bins. Also by setting the `normed` parameter to 1 the counts are normalized to form a probability density.

```


7
8 import numpy as np
9 import matplotlib.pyplot as plt
10
11 mu, sigma = 100, 15
12 x = mu + sigma * np.random.randn(1000)
13 n, bins, patches = plt.hist(x, 10, \
14                             normed=1, \
15                             facecolor='g')
16 plt.xlabel('Frequency')
17 plt.ylabel('Probability')
18 plt.title('Histogram Example')
19 plt.text(40, .028, 'mean=100 std.dev.=15')
20 plt.axis([40, 160, 0, 0.03])
21 plt.grid(True)
22 plt.show()

```



Next lets look at the scatter plot. The scatter object takes two sequence objects, such as arrays, of the same length and optional parameters to denote color and style attributes.

```
7
8 import numpy as np
9 import matplotlib.pyplot as plt
10
11 N = 100
12 x = np.random.rand(N)
13 y = np.random.rand(N)
14 colors = np.random.rand(N)
15 #colors=('r', 'b', 'g')
16 area = np.pi * (10 * np.random.rand(N))**2 # 0
17 plt.scatter(x, y, s=area, c=colors, alpha=0.5)
18 plt.show()
19
20
```



The Pandas library builds on NumPy by introducing several useful data structures and functionalities to read and process data. Pandas is a great tool for general data munging. It easily handles common tasks such as dealing with missing data, manipulating shapes and sizes, converting between data formats and structures, and importing data from different sources.

The main data structures introduced by Pandas are:

- Series
- The DataFrame
- Panel




The DataFrame is probably the most widely used. It is a two-dimensional structure that is effectively a table created from either a NumPy array, lists, dicts, or series. You can also create a DataFrame by reading from a file.

The Pandas library builds on NumPy by introducing several useful data structures and functionalities to read and process data. Pandas is a great tool for general data munging. It easily handles common tasks such as dealing with missing data, manipulating shapes and sizes, converting between data formats and structures, and importing data from different sources.

The main data structures introduced by Pandas are:

- Series
- The DataFrame
- Panel

The DataFrame is probably the most widely used. It is a two-dimensional structure that is effectively a table created from either a NumPy array, lists, dicts, or series. You can also create a DataFrame by reading from a file.



Lets consider Pandas by working through some data.


An interesting example is the historical weather observations from the Hobart weather station in Tasmania and consider the task of “*discovering how the daily maximum temperature has changed over time*”.




First lets create a DataFrame from the data
And display the first few lines of it

```
import pandas as pd

df = pd.read_csv('C:/Users/Jake/sampleData.csv')

df.head()
```








The product code and the station number are the same for each row and that this information is superfluous. Also, the days of accumulated maximum temperature are not needed for our purpose, so we will delete them as well:

```
In [47]: runfile('C:/Users/Jake/lane.py', wdir='C:/Users/Jake')
```

	Product code	Bureau of Meteorology station number	Year	Month	Day	\
0	IDCJAC0010	94029	1882	1	1	
1	IDCJAC0010	94029	1882	1	2	
2	IDCJAC0010	94029	1882	1	3	
3	IDCJAC0010	94029	1882	1	4	
4	IDCJAC0010	94029	1882	1	5	

	Maximum temperature (Degree C)	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

Days of accumulation of maximum temperature Quality



The product code and the station number are the same for each row and that this information is superfluous. Also, the days of accumulated maximum temperature are not needed for our purpose, so we will delete them as well:

```




/
8 import pandas as pd
9
10 df = pd.read_csv('C:/Users/Jake/sampleData.csv')
11
12 del df['Bureau of Meteorology station number']
13 del df['Product code']
14 del df['Days of accumulation of maximum temperature']
15
16 df.head()
17
18
19
20
21

```

```

Console 1/A
Out[54]:
      Year  Month  Day  Maximum temperature
(Degree C) Quality
0  1882      1    1
NaN      NaN
1  1882      1    2
NaN      NaN
2  1882      1    3
NaN      NaN
3  1882      1    4
NaN      NaN
4  1882      1    5
NaN      NaN

```



Let's make our data a little easier to read by shorting the column labels:

```
8 df=df.rename(columns={'Maximum temperature (Degree C)': 'maxtemp'})
```

```
13
14 del df['Bureau of Meteorology station number']
15 del df['Product code']
16 del df['Days of accumulation of maximum temperature']
17
18 df=df.rename(columns={'Maximum temperature (Degree C)': 'maxtemp'})
19
20 df.head()
21
22
```

```
wdl1 = C:/Users/Jake
```

```
In [2]: df.head()
```

```
Out[2]:
```




	Year	Month	Day	maxtemp	Quality
0	1882	1	1	NaN	NaN
1	1882	1	2	NaN	NaN
2	1882	1	3	NaN	NaN
3	1882	1	4	NaN	NaN
4	1882	1	5	NaN	NaN

We are only interested in data that is of high quality, so we include only records that have a Y in the quality column:

```
17
18 df=df.rename(columns={'Maximum temperature (Degree C)': 'maxtemp'})
19
20 df=df[(df.Quality=='Y')]
21
22
23
24
```

```
Out[5]:
```

	Year	Month	Day	maxtemp	Quality
91	1882	4	2	15.6	Y
92	1882	4	3	14.4	Y
93	1882	4	4	12.8	Y
94	1882	4	5	10.0	Y
95	1882	4	6	10.6	Y






We can get a statistical summary of our data using the describe() command

```
In [6]: df.describe()
```


```
Out[6]:
```

	Year	Month	Day	maxtemp
count	44250.000000	44250.000000	44250.000000	44250.000000
mean	1952.207503	6.536339	15.734870	16.929941
std	38.212270	3.446311	8.802089	5.030362
min	1882.000000	1.000000	1.000000	4.300000
25%	1924.000000	4.000000	8.000000	13.300000
50%	1954.000000	7.000000	16.000000	16.400000
75%	1985.000000	10.000000	23.000000	20.000000
max	2015.000000	12.000000	31.000000	41.800000



Lets look at our data in Matplotlib

```
24 import matplotlib.pyplot as plt
25 plt.plot(df.Year, df.maxtemp)
26
```



SciPy (pronounced sigh pi) adds a layer to NumPy that wraps common scientific and statistical applications on top of the more purely mathematical constructs of NumPy.

SciPy provides higher-level functions for manipulating and visualizing data, and it is especially useful when using Python interactively. SciPy is organized into sub-packages covering different scientific computing applications.

A list of the packages most relevant to ML and their functions appear as follows:


Package	Description
<code>cluster</code>	This contains two sub-packages: <code>cluster.vq</code> for K-means clustering and vector quantization. <code>cluster.hierarchy</code> for hierarchical and agglomerative clustering, which is useful for distance matrices, calculating statistics on clusters, as well as visualizing clusters with dendrograms.
<code>constants</code>	These are physical and mathematical constants such as π and e .
<code>integrate</code>	These are differential equation solvers
<code>interpolate</code>	These are interpolation functions for creating new data points within a range of known points.
<code>io</code>	This refers to input and output functions for creating string, binary, or raw data streams, and reading and writing to and from files.
<code>optimize</code>	This refers to optimizing and finding roots.
<code>linalg</code>	This refers to linear algebra routines such as basic matrix calculations, solving linear systems, finding determinants and norms, and decomposition.
<code>ndimage</code>	This is N-dimensional image processing.
<code>odr</code>	This is orthogonal distance regression.
<code>stats</code>	This refers to statistical distributions and functions.

This includes algorithms for the most common machine learning tasks, such as classification, regression, clustering, dimensionality reduction, model selection, and preprocessing.

Scikit-learn comes with several real-world data sets for us to practice with. Let's take a look at one of these—the Iris data set:

```
7
8 from sklearn import datasets
9
10 iris = datasets.load_iris()
11
12 iris_X = iris.data
13 iris_y = iris.target
14
15 print(iris_X.shape)
16
17 |
```

)
(150, 4)
>>>



The data set contains 150 samples of three types of irises (Setosa, Versicolor, and Virginica), each with four features.

We can see that the four attributes, or features, are sepal width, sepal length, petal length, and petal width in centimeters. Each sample is associated with one of three classes. **Setosa**, **Versicolor**, and **Virginica**. These are represented by 0, 1, and 2 respectively.

Let's look at a simple classification problem using this data. We want to predict the type of iris based on its features: the length and width of its sepal and petals. Typically, scikit-learn uses estimators to implement a $\text{fit}(X, y)$ method and for training a classifier, and a $\text{predict}(X)$ method that if given unlabeled observations, X , returns the predicted labels, y . The $\text{fit}()$ and $\text{predict}()$ methods usually take a 2D array-like object. {[more on this next semester](#)}